

- Score Normalization (Min-Max, RRF)
- Dynamic Weighting (λ based on query)
- Optional Cross-Encoder Reranking



- PDF Storage (Amazon S3)
- Pre-signed URLs
- Metadata, previews



- Logging & Monitoring Stack
- Prometheus, Grafana, ELK
- AWS CloudWatch



- Indexing Pipeline (Lambda)
- Triggered on S3 Upload
- Text Extraction & Embedding
- Update OpenSearch & Pinecone
- Refresh MongoDB mappings

Key Components & How They Address Issues

1. User Authentication & RBAC

- **What:** FastAPI validates JWT tokens and extracts user roles.
- **Why:** Ensures that only authorized users access PDFs according to permissions defined in MongoDB.

2. Pre-filtering Layer

- **What:** Queries MongoDB for user-PDF mappings and caches results in Redis.
- **Why:** Reduces database load and latency on frequent role-to-PDF lookups.

3. Query Caching

- **What:** Caches frequent search queries in Redis (with a short TTL).
- **Why:** Improves performance by avoiding redundant search calls.

4. Hybrid Search Engine

- **What:** Executes BM25 search via OpenSearch and semantic search via Pinecone in parallel (using async calls).
- **Why:** Combines exact keyword matching with contextual similarity while minimizing latency.

5. Ranking & Fusion Service

- **What:** Normalizes scores, applies dynamic weighting (or RRF), and optionally uses a cross-encoder for reranking top results.
- **Why:** Ensures the final ranked list reflects both relevance and query intent, adapting to different query types.

6. PDF Storage (Amazon S3)

- **What:** Stores PDFs securely and provides pre-signed URLs for controlled access.
- **Why:** Protects PDF access while ensuring that metadata and previews are returned to aid user selection.

7. Indexing Pipeline (Lambda)

- **What:** Automatically triggers on new PDF uploads to extract text, generate embeddings, and update search indices.
- **Why:** Keeps your search indices up-to-date without manual re-indexing.

8. Error Handling & Fallbacks

- **What:** FastAPI's try-catch blocks, fallback to direct MongoDB queries if Redis fails, partial result returns if a search component is down.
- **Why:** Enhances reliability and ensures graceful degradation in case of component failures.

9. Scalability Management & Observability

- **What: Monitors system performance with Prometheus, Grafana, ELK, and AWS CloudWatch; scales OpenSearch/Pinecone as needed.**
- **Why: Proactively tracks and manages performance bottlenecks and resource demands.**

File structure

```

├─ app/
|   ├── __init__.py
|   ├── main.py          # FastAPI app entry point with exception handlers
|   ├── config.py        # Environment configurations
|   ├── auth.py          # JWT authentication and RBAC utilities
|   ├── exceptions/      # Custom exceptions
|   |   ├── __init__.py
|   |   └─ custom_errors.py
|   ├── models/
|   |   ├── __init__.py
|   |   ├── user.py      # Pydantic models for users
|   |   └─ pdf.py        # Models for PDF metadata and results
|   ├── routes/
|   |   ├── __init__.py
|   |   └─ search.py     # API endpoints
|   ├── services/
|   |   ├── __init__.py
|   |   ├── pdf_filter.py # Pre-filtering with Redis/MongoDB
|   |   ├── hybrid_search.py # Parallel BM25 & vector search with snippets
|   |   ├── ranking.py   # Ranking and fusion logic
|   |   └─ s3_storage.py  # S3 URL generation
|   ├── workers/         # Background tasks
|   |   ├── __init__.py

```

- | | |─ tasks.py # Cross-encoder reranking, etc.
- | | |─ indexer.py # Incremental indexing logic
- | |─ utils/
- | | |─ __init__.py
- | | |─ helpers.py # General utilities
- | | |─ query_cache.py # Query result caching
- | | |─ logging.py # Structured logging
- | | |─ metrics.py # Prometheus metrics
- | |─ lambda/ # (Optional) AWS Lambda for indexing
- | | |─ indexer_lambda.py # Lambda handler for S3 events
- | |─ requirements.txt
- |─ tests/ # Unit and integration tests
- | |─ __init__.py
- | |─ test_auth.py
- | |─ test_search.py
- | |─ test_services.py
- |─ .apprunner.yaml # AWS App Runner config
- |─ Dockerfile # Docker build instructions
- |─ requirements.txt # Dependencies
- |─ README.md # Documentation

project/

- |─ app/
- | |─ __init__.py
- | |─ main.py
- | |─ config.py



- | |─ auth.py
- | |─ models/
 - | |─ user.py
 - | |─ pdf.py
- | |─ routes/
 - | |─ search.py
- | |─ services/
 - | |─ pdf_filter.py
 - | |─ hybrid_search.py
 - | |─ ranking.py
 - | |─ s3_storage.py
- | |─ utils/
 - | |─ helpers.py
 - | |─ query_cache.py
- | |─ exceptions/
 - | |─ custom_errors.py
- | |─ workers/
 - | |─ tasks.py
 - | |─ indexer.py
- |─ lambda/
 - |─ indexer_lambda.py
- |─ tests/
 - |─ test_auth.py
- |─ Dockerfile
- |─ .apprunner.yaml
- |─ requirements.txt
- |─ .env # Not included here; user must create this
- |─ README.md

Main Functions Defined In our project:-



1. User Uploads PDF: Done

-  File Upload Component
-  Uploads PDF to AWS S3


2. Processing the PDF: Done

-  Extract text from the PDF
-  Preprocess (clean, remove unnecessary characters)


3. Chunking the Text: Done

-  Divide text into smaller chunks
-  Maintain contextual meaning

4. Embedding the Chunks: Done

-  Convert chunks into vector embeddings using an NLP model Sentence Transformer

5. Storing the Chunks: half Done

-  Store the vector embeddings in Pinecone for retrieval

- Store Metadata in MongoDB

6. Search & Retrieval

- 🔍 Query processing using OpenSearch
- 📊 Find similar text chunks from Pinecone
- 🏆 Return best-matching results

Authentication and Other feature in our project:

1. User Registration: Done ✅

- 📄 User enters details (username, email, password, etc.)
- 📧 System sends OTP to user's email



2. OTP Verification: Done ✅

- 🔑 User enters OTP received via email
- ✅ System verifies OTP
- 🎉 User is successfully registered

3. User Login: Done ✅



- 🗝️ User enters credentials
- 🔑 System verifies credentials
- 🎫 System generates authentication token

4. Upload Route Access: **Done**




-  Authenticated users can now access the **Upload PDF** route
-  User uploads PDF (leads to PDF processing flow)

Flow Structure for Profile Page (CRUD Operations)



1. User Authentication & Profile Access: **Done**

-  User logs in (Token-based authentication)
-  User navigates to the profile page



2. Update User Details (PUT Request): **Done**


-  User modifies profile details (name, email, password, etc.)
-  System updates change in the database
-  Confirmation message is shown

3. Fetch Uploaded PDFs (GET Request):





-  System retrieves all PDFs uploaded by the user
-  User sees a list of uploaded PDFs

4. Delete a PDF (DELETE Request):




-  User selects a PDF to delete
-  System removes the PDF from AWS S3 & database

-  Confirmation message is shown

5. Delete Account (DELETE Request): Done



-  User requests account deletion
-  System verifies user action (optional re-authentication)
-  Deletes all user data (profile + PDFs)
-  Logs out user and redirects to the home page

6. Logout (Token Expiry) Done

-  User clicks logout
-  System clears authentication token
-  Redirects to login page

Folder Structure and their Works:

bucket/ (Manages AWS & Pinecone Initialization)

-  **AWSBucket.py** – Initializes AWS S3 for file storage
 - Configures AWS SDK
 - Handles file uploads & deletions
-  **PineconeBucket.py** – Initializes Pinecone for vector storage
 - Connects to Pinecone
 - Manages vector embeddings storage & retrieval

controller/ (Handles Routing, No Business Logic)

This folder defines API endpoints and delegates business logic to injected services.

AccountController.py

◆ **Manages User Authentication Routes**

- POST /register → Calls service to register a new user & send OTP.
- POST /login → Calls service to authenticate user & return token.

✓ **Dependencies Passed to Services:**

- Auth service (Handles authentication & token generation)

UploadController.py

◆ **Manages File Upload & Processing Routes**

- POST /upload → Calls service to handle PDF upload to AWS S3.
- GET /fetch-files → Calls service to retrieve user's uploaded PDFs.

✓ **Dependencies Passed to Services:**

- `s3_service` (Uploads & deletes files from AWS S3)
- `text_service` (Extracts & processes text from PDFs)
- `embedding_service` (Embeds chunks & stores in Pinecone)
- `pinecone_service` (Manages vector storage & retrieval)

db/ (Manages Database Connection & GridFS Integration)

connection.py

◆ **Initializes database connection & GridFS for file storage**

- Connects to MongoDB using **Mongo Client**
- Creates a user's collection for storing user data
- Initializes **GridFS** for handling large file uploads (PDFs)
- Ensures the database connection is accessible across the app

✓ **Key Components:**

- `client = MongoClient (MONGO_URI) →`
Connects to MongoDB
- `db = client["your_database_name"] →` Selects database
- `users_collection = db["users"] →` Defines user collection
- `grid_fs = GridFS(db) →` Initializes GridFS for file storage

collection/ (Defines Database Models)

userCollection.py

◆ **Defines User Registration Schema**

- Specifies the **structure of user documents** in MongoDB
- Ensures fields like **email**, **password_hash**, **created_at**, etc.
- Implements any pre-processing (e.g., hashing passwords before storing)

diInjector/ (Manages Dependency Injection for Services)

diExtension.py

◆ **Injects Core Services into the Application**

- **Registers services globally** so controllers can access them
- Ensures **loose coupling** (controllers don't directly instantiate services)
- Helps with **scalability & testing** (easily replaceable service implementations)

dtos/ (Data Transfer Objects)

- Structures and validates incoming user data before processing.
- Ensures correct format for **register**, **login**, **upload**, etc.
- Example: **LoginDTO**, **RegisterDTO**, **UploadDTO**.

schema/ (Validation Layer)

- Defines schemas to validate user input before storing in the database.
- Ensures that required fields are provided and formatted correctly.
- Example:
 - `register_schema.py`: Validates user registration details.
 - `upload_schema.py`: Ensures correct file format and size.

3 **routes/ (API Endpoints)**

- Defines all API routes for user interaction.
- Forwards requests to the **controller**.
- Example routes:
 - `POST /register` → Calls `AccountController`
 - `POST /upload` → Calls `UploadController`

service/ (Business Logic Layer)

- Implements actual **business logic** for register, login, and upload.
- Calls **database operations** and handles **tokens, file uploads, etc.**
- Example:
 - `account_service.py`: Handles user creation, password hashing, and authentication.
 - `upload_service.py`: Processes PDFs, extracts text, creates embeddings.

JWTConfig/ (Authentication & Token Management)

- Generates and verifies authentication tokens for users.
- Manages token expiration and refresh logic.

middleware/ (Access Control & Security)

- Restricts unauthorized users from accessing certain routes.
- Ensures only authenticated users with **client ID** can perform actions.

utils/ (Helper & Utility Functions)

- Contains reusable helper functions like:
 - **Embedding chunks of text** (for Pinecone storage).
 - **Sending emails** (for OTP verification, password reset, etc.).
 - **File handling** (e.g., processing PDFs before uploading).

Folder Responsibilities

Folder	Responsibility
bucket/	Initializes AWS S3 & Pinecone for filestorage
controller/	Defines API endpoints but no business logic
db/	Manages MongoDB connection (GridFS for PDF storage)

collection/	Defines user schema & database models
diInjector/	Injects services into controllers using Dependency Injection
dtos/	Defines Data Transfer Objects (DTOs) for structured API requests
enum/	Stores predefined constants (user roles, status codes, etc.)
exception/	Handles custom error handling & predefined exceptions
helper/	Stores helper functions (password hashing, JWT handling)
interface/	Defines service interfaces for register, login, upload, etc.
JWTconfig/	Manages JWT configuration (secret key, token expiry)
lambda/	AWS Lambda functions (if applicable)
middleware/	Restricts user access based on <code>client_id</code>
model/	Defines expected input/output data models
routes/	Registers all API routes
schema/	Handles schema validation (registration, login, upload)
service/	Implements business logic (register, login, upload, embedding)

template/	Stores email templates (OTP, welcome email, etc.)
utils/	Stores utility functions (embedding, email sending, file processing)